

The Multi-Layer Perceptron (MLP)

a short introduction

Cornelius Weber
Knowledge Technology Research Group
Informatik, Universität Hamburg
Email: weber@informatik.uni-hamburg.de

October 7, 2014

1-Layer Perceptron = a layer of neurons

activate a neuron:

$$h_j^{out} = \sum_n w_{jn} s_n^{in} = \vec{w}_j \cdot \vec{s}^{in}$$

→ dot product (scalar product)
between weight vector and input vector

activate all neurons:

$$\vec{h}^{out} = W \vec{s}^{in}$$


→ matrix product with weight matrix

in Python: `h_out = numpy.dot(W,s_in)`

in C: two nested for-loops

(outer loop over output neurons, inner loop does scalar product)

transfer function applied, e.g. $s_j^{out} = \tanh(h_j^{out})$



Feedforward Activation of the MLP

inner activation of neuron i in layer λ :


$$h_i^\lambda(\vec{w}_i^{\lambda\lambda-1}) = \sum_j w_{ij}^{\lambda\lambda-1} s_j^{\lambda-1} + \underbrace{b_i^\lambda}_{\text{bias}}$$

neuron output:


$$s_i^\lambda(h_i) = \varphi_i^\lambda(h_i^\lambda) \quad \stackrel{\text{e.g.}}{=} \tanh(h_i^\lambda)$$

in vector notation:

$$\vec{s}^\lambda = \varphi^\lambda(\underbrace{W^{\lambda\lambda-1} \vec{s}^{\lambda-1}}_{\vec{h}^{\lambda-1}})$$



Feedforward Activation of the MLP



$$\vec{s}^{out} = \varphi^{out}(W^{out} \underbrace{\varphi^{hid}(W^{hid} \vec{s}^{in})}_{\vec{s}^{hid}})$$

Feedforward Activation of the MLP

inner activation of neuron i in layer λ :

$$h_i^\lambda(\vec{w}_i^{\lambda, \lambda-1}) = \sum_j w_{ij}^{\lambda, \lambda-1} s_j^{\lambda-1} + \underbrace{b_i^\lambda}_{\text{bias}}$$

neuron output:


$$s_i^\lambda(h_i) = \varphi_i^\lambda(h_i^\lambda) \quad \text{e.g.} \quad \tanh(h_i^\lambda)$$

local error on output layer:

$$e_i(s_i) = \textcolor{red}{s}_i^{\text{Teach}} - s_i^{\text{out}}$$

total error (cost function):

$$E(\vec{e}) = \frac{1}{2} \sum^{data} \sum_i e_i^2$$



Dependencies: $E(e(s(h(w))))$. Chain rule \Rightarrow derivatives for training

Possible Transfer Functions

- ▶ linear transfer function — typically used for output units

$$\varphi_i = h_i = \sum_j w_{ij} s_j^{in} + b_i \quad \varphi'_i = \frac{\partial \varphi_i}{\partial h_i} = 1$$

- ▶ logistic/sigmoid function — similar shape as tanh but in $]0, 1[$

$$\varphi_i = \frac{1}{1 + e^{-h_i}} = \frac{1}{1 + e^{-\sum_j w_{ij} s_j^{in} + b_i}} \quad \varphi'_i = \varphi_i \cdot (1 - \varphi_i)$$

- ▶ hyperbolic tangent

$$\varphi_i = \tanh(h_i) \quad \varphi'_i = 1 - (\varphi_i)^2$$

- ▶ Radial Basis Functions — only used for \vec{s}^{in} on input layer

$$\varphi_i = e^{-\frac{\sum_j (s_j^{in} - w_{ij})^2}{2\sigma_i^2}}$$

- ▶ max-like operation (Riesenhuber & Poggio's "HMAX model")

$$\varphi_i = \max_j w_{ij} s_j^{in} \quad \textit{not differentiable}$$

Non-Local “Layer” Transfer Functions

$$\text{let } h_i = \sum_j w_{ij} s_j^{\text{in}} + b_i$$

- ▶ softmax (inverse temperature β)

$$\varphi_i = \frac{e^{\beta h_i}}{\sum_k e^{\beta h_k}}$$

- ▶ winner-take-all (not differentiable)

$$\varphi_{i^*} = \begin{cases} 1, & h_{i^*} > h_k \ \forall k \neq i^* \\ 0, & \text{else} \end{cases}$$

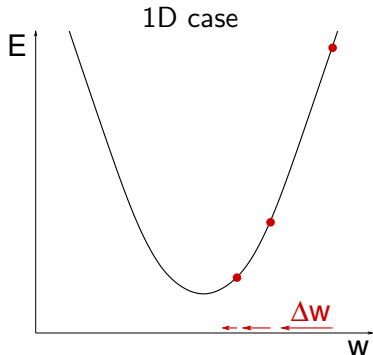
- ▶ competitive topographic (given i^* is winning node)

$$\varphi_i = e^{-\frac{(i-i^*)^2}{\sigma^2}}$$


note: non-local transfer functions lead to non-local update rules from backpropagation

Gradient Descent

$$\Delta w = -\frac{\partial E}{\partial w}$$



for quadratic functions,
update size scales very well



update size in one direction
does not match update size in
another

Natural Gradient

$$\Delta w = - G^{-1} \frac{\partial E}{\partial w}$$


Compensates for “distorted” coordinates. For affine distortions:

$$G = J^T J = \begin{pmatrix} \frac{\partial^2 E}{\partial w_1^2} & \frac{\partial^2 E}{\partial w_1 \partial w_2} & \cdots \\ \frac{\partial^2 E}{\partial w_1 \partial w_2} & \frac{\partial^2 E}{\partial w_2^2} & \cdots \\ \cdots & \cdots & \cdots \end{pmatrix} = H \text{ (Hessian)}$$

$$J = \left(\frac{\partial E}{\partial w_1}, \frac{\partial E}{\partial w_2}, \dots \right) = \text{Jacobian (here, a vector)}$$

Approximation, neglecting off-diagonals:

$$G^{-1} \approx \begin{pmatrix} \frac{1}{\frac{\partial^2 E}{\partial w_1^2}} & 0 \\ 0 & \frac{1}{\frac{\partial^2 E}{\partial w_2^2}} \end{pmatrix}$$



Error Backpropagation: Recursive δ

last layer (using index i):

$$\begin{aligned}\frac{\partial E}{\partial w_{ij}} &= \frac{\partial E}{\partial e_i} \frac{\partial e_i}{\partial s_i} \frac{\partial s_i}{\partial h_i} \frac{\partial h_i}{\partial w_{ij}} \\ &= \underbrace{e_i - 1}_{\delta_i} \varphi'(h_i) s_j\end{aligned}$$

2nd-last layer (using index j):

$$\begin{aligned}\frac{\partial E}{\partial w_{jk}} &= \underbrace{\sum_i \frac{\partial E}{\partial e_i} \frac{\partial e_i}{\partial s_j}}_{\delta_i \text{ (compare with above)}} \cdot \frac{\partial s_j}{\partial w_{jk}} \\ &= \underbrace{\sum_i \underbrace{\frac{\partial e_i}{\partial h_i}}_{-\varphi'_i(h_i)} \underbrace{\frac{\partial h_i}{\partial s_j}}_{w_{ij}}}_{\delta_i \text{ (compare with above)}} \cdot \underbrace{\frac{\partial s_j}{\partial h_j} \frac{\partial h_j}{\partial w_{jk}}}_{\varphi'_j(h_j) s_k} \\ &= \delta_j \text{ (defined as } \delta_j = \varphi'_j \sum_i \delta_i w_{ij} \text{)}\end{aligned}$$

One MLP Learning Step

Given input activities s_j^0 on layer $\lambda = 0$.

Forward propagation of activities to hidden- and output layers:

$$s_i^\lambda = \varphi\left(\sum_j w_{ij}^{\lambda, \lambda-1} s_j^{\lambda-1} + b_i^\lambda\right)$$

Error on output layer: *(see next 5 slides)*

$$\delta_i^{\text{out}} = \varphi'_i(h_i^{\text{out}}) \cdot (s_i^{\text{Teach}} - s_i^{\text{out}}) \stackrel{\text{usually}}{=} s_i^{\text{Teach}} - s_i^{\text{out}}$$

Error backpropagation to all hidden layers:

$$\delta_j^\lambda = \varphi'_j(h_j^\lambda) \cdot \sum_i \delta_i^{\lambda+1} w_{ij}^{\lambda+1, \lambda}$$

Learning rule, with small learning rate ϵ :

$$\left. \begin{aligned} \Delta w_{ij}^{\lambda, \lambda-1} &= -\epsilon \frac{\partial E}{\partial w_{ij}^{\lambda, \lambda-1}} = \epsilon \delta_i^\lambda s_j^{\lambda-1} \\ \Delta b_i^\lambda &= -\epsilon \frac{\partial E}{\partial b_i^\lambda} = \epsilon \delta_i^\lambda \end{aligned} \right\}$$

Many learning steps with different data samples must be repeated ...

Interpretation of the Squared Error

likelihood of the data s^{Teach} being generated
given input s^{in} and model parameters W

$$P(s^{\text{Teach}}|W, s^{\text{in}}) \approx e^{-\frac{1}{2}(s^{\text{Teach}} - s^{\text{out}})^2} = e^{-E}$$

assumption: Gaussian noise on output

- ▶ $E = -\ln P$
- ▶ $P^a P^b = e^{-E^a} e^{-E^b} = e^{-E^a - E^b}$
multiplying probabilities \Leftrightarrow adding cost terms

note: Gaussian noise unbounded \Rightarrow unbounded output range
assumed \Rightarrow linear transfer function on output units appropriate,
but not sigmoid function

Derivative for Square Error & Linear Function

given: $s^{out} := \varphi(wx) = wx$ linear transfer function

$$\begin{aligned}-\frac{\partial E}{\partial w} &= -\frac{\partial}{\partial w} \left(\frac{1}{2} (s^{\text{Teach}} - s^{out})^2 \right) \\ &= (s^{\text{Teach}} - s^{out}) \frac{\partial s^{out}}{\partial w} \\ &= (s^{\text{Teach}} - s^{out}) x\end{aligned}$$

→ squared error & linear function \Rightarrow simple learning rule

Cross-Entropy Error

outputs considered as class probabilities:

$$\begin{aligned}P(s^{\text{Teach}}|W, s^{\text{in}}) &= \begin{cases} s^{\text{out}}, & s^{\text{Teach}} = 1 \quad \text{it's this class} \\ 1 - s^{\text{out}}, & s^{\text{Teach}} = 0 \quad \text{not this class} \end{cases} \\ &= (s^{\text{out}})^{s^{\text{Teach}}} \cdot (1 - s^{\text{out}})^{(1-s^{\text{Teach}})}\end{aligned}$$

cross-entropy error:

$$\begin{aligned}E &= -\ln P \\ &= -s^{\text{Teach}} \ln s^{\text{out}} - (1 - s^{\text{Teach}}) \ln(1 - s^{\text{out}})\end{aligned}$$

note: for 2 classes, sigmoid/logistic transfer function is natural

www.willamette.edu/~gorr/classes/cs449/classify.html

Derivative for Cross Entropy & Logistic Function

given: $s^{out} := \varphi(wx) = \frac{1}{1+e^{-wx}}$ logistic transfer function

for the logistic function it is: $\frac{\partial s^{out}}{\partial w} = x s^{out} (1-s^{out})$

$$\begin{aligned}-\frac{\partial E}{\partial w} &= -\frac{\partial}{\partial w} \left(-s^{\text{Teach}} \ln s^{out} - (1-s^{\text{Teach}}) \ln(1-s^{out}) \right) \\&= s^{\text{Teach}} \frac{1}{s^{out}} \frac{\partial s^{out}}{\partial w} - (1-s^{\text{Teach}}) \frac{1}{1-s^{out}} \frac{\partial s^{out}}{\partial w} \\&= s^{\text{Teach}} \frac{1}{s^{out}} x s^{out} (1-s^{out}) - (1-s^{\text{Teach}}) \frac{1}{1-s^{out}} x s^{out} (1-s^{out}) \\&= s^{\text{Teach}} x (1-s^{out}) - (1-s^{\text{Teach}}) x s^{out} \\&= \begin{cases} s^{\text{Teach}} x (1-s^{out}), & s^{\text{Teach}} = 1 \\ -(1-s^{\text{Teach}}) x s^{out}, & s^{\text{Teach}} = 0 \end{cases} \\&= x (s^{\text{Teach}} - s^{out})\end{aligned}$$

→ cross-entropy & logistic function \Rightarrow simple learning rule

Derivative for Cross Entropy & Softmax Function

given: $s_k^{out} := \varphi(\{w_{k'}\}, x) = \frac{e^{w_k x}}{\sum_{k'} e^{w_{k'} x}}$ softmax transfer function
useful for 1-of-K classification; output units' activations dependent
for softmax it is: $\frac{\partial s_k^{out}}{\partial w_k} = x s_k^{out} (1 - s_k^{out})$ (as for logistic)


$$\begin{aligned} -\frac{\partial E}{\partial w_k} &= -\frac{\partial}{\partial w_k} \left(-\sum_k s_k^{Teach} \ln s_k^{out} \right) \\ &= \dots \\ &= x (s_k^{Teach} - s_k^{out}) \end{aligned}$$

→ cross-entropy & softmax function \Rightarrow simple learning rule

note, for 2 classes ($K = 2$), softmax becomes logistic function:

$$\begin{aligned} s_1^{out} &= \frac{e^{w_1 x}}{e^{w_1 x} + e^{w_2 x}} = \frac{1}{1 + e^h} & \text{with } h = (w_2 - w_1)x \\ s_2^{out} &= \frac{e^{w_2 x}}{e^{w_1 x} + e^{w_2 x}} = \frac{1}{1 + e^{-h}} \end{aligned}$$

Error Backpropagation – Comments



- ▶ non-linearity on middle layer important
- ▶ more than linear separation – complex transformations possible
- ▶ 3-layer network can in theory represent any input-output function
- ▶ more layers possible → “deep learning” (requires tricks, e.g. initial unsupervised learning, convolutional kernel, max-pooling, ...)
- ▶ backpropagation considered unbiological
- ▶ overfitting possible

Krizhevsky Sutskever Hinton. ImageNet Classification with Deep Convolutional Neural Networks. NIPS 2012

Ciresan Meier .. Schmidhuber. Multi-Column Deep Neural Network for Traffic Sign Classification. Neur Netw 2013

Early Stopping

- ▶ use network with less parameters
- ▶ stop when test error increases

